

场景 A • 宿舍公平助手 — 完整聊天痕迹

SafeBARS Ethical Mirror 真实 session 重放 (StudentTest/study_bundle.json)

⚠ 诚实说明：本页 100% 取自真人测试捕获数据。研究者**做了决策** (add_safeguard / consult_stakeholders)，但**计划文本具体改了什么未被记录** (台账: "No concise changed passage was recorded") → 「找得到」证据强, 「纠正」证据弱。

Step 1 • 研究者输入研究计划

Research area and context I want to explore fairness in shared dorm living; roommates often argue about who cleans and who free-rides on chores. Intended change Gently remind the free-riding roommate to restore fair chore-sharing and reduce built-up resentment. Direct users and encounter University students living in dorms, including the user and their roommates. AI role and decision authority AI reads smart-plug electricity data and group-chat tone to judge who is free-riding, then drafts a private message the user can send to the roommate. Data and materials Roommates electricity usage logs, group-chat text, and possibly camera or door-sensor presence records. Sensitive-data necessity and alternative Roommates may feel watched and labelled as lazy; a solvable conflict turns into algorithmic blame; there is also a privacy-leak risk. People affected without direct use The tool only gives gentle prompts, never public naming; and it asks co-residents for consent before collecting any data. Researcher-authored value commitment If data shows it is deepening conflict or being used to bully a roommate, stop it and redesign. Pause or redesign condition If reports show the tool increases roommate conflict, gets used to shame or harass someone, or a data breach exposes private chat or presence records, I would pause distribution and redesign: switch to opt-in only, drop any camera or door-sensor data, and add a human-review step before any message is sent.

Step 2-4 • AI 通过 9 个文献镜头分析并浮现盲点

AI: 我分析了你的计划, 发现如下伦理问题 (徽章=证据覆盖状态):

Lifecycle integration **已论证**

The plan names pause/redesign conditions tied to specific observed outcomes (e.g., increased conflict, shaming, data breach) but does not map these to explicit pre-deployment, during-deployment, or post-deployment checkpoints.

↳ 自问: Where will unintended consequences be reconsidered before, during, and after design, evaluation, and deployment?

Benefit-harm assumptions **已论证**

The plan names benefits (restoring fair chore-sharing, reducing resentment) and harms (algorithmic blame, privacy leaks, deepening conflict), and links them to mechanisms like tone analysis misjudgment or camera data exposure.

↳ 自问: Which promised benefits and harm assumptions could fail, for whom, and through what mechanism?

Affected parties and distribution **已声明**

The plan mentions 'co-residents' and 'roommates' but only frames them as consent-givers—not as potentially burdened parties whose autonomy, reputation, or mental health may be affected without their active participation in design.

↳ 自问: Who may benefit, bear burdens, lose access, or be affected without being a direct user or study participant?

Downstream use, misuse, and scale **缺失**

The plan does not consider how outputs (e.g., drafted messages, inferred 'free-rider' labels) could be repurposed outside the intended context—e.g., shared with landlords, cited in conduct hearings, or integrated into university disciplinary systems.

↳ 自问: How could the research output be repurposed, misused, combined with other systems, or deployed at a scale beyond the study?

Perspective diversity and participation **已声明**

The plan states consent is sought but gives no method, timing, or mechanism for incorporating roommate perspectives *before* design finalization or *during* iteration—only passive permission, not co-shaping.

↳ 自问: Whose situated knowledge is needed, and when will real affected people be able to challenge the team's assumptions?

Responsibility, oversight, and contestability **已声明**

The plan identifies the researcher as accountable for pausing/redesigning but offers no mechanism for affected roommates to understand, challenge, or appeal an AI-generated message or label.

↳ 自问: Who remains accountable for AI-assisted decisions, and how can an affected person understand, challenge, override, or appeal them?

Prior cases and emerging-capability horizon **缺失**

The plan cites no prior cases (e.g., workplace wellness monitoring, academic integrity tools) or emerging capabilities (e.g., multimodal LLMs inferring intent from partial data) that inform its risk assessment or trigger points.

↳ 自问: Which prior incidents and current AI capabilities challenge the plan, and what new capability would trigger reassessment?

Mitigation and design-change commitment **已联动**

The plan binds three specific harms (increased conflict, shaming/harassment, data breach) to concrete design changes: opt-in-only mode, dropping camera/door-sensor data, and adding human review before message delivery.

↳ 自问: Which concrete feature, scope, method, safeguard, or stopping rule will change if the concern is credible?

Monitoring, learning, and redress **已声明**

The plan references 'reports' as triggers but specifies no indicators, cadence, complaint channel, redress process, or accountability for follow-up—only a reactive pause condition.

↳ 自问: How will the team detect harm, receive complaints, correct errors, support remedy, and learn after release or study completion?

Step 5 · AI 发现承诺—失效张力边 + 研究者回应

AI：我还发现 2 条张力，建议你处理：

⚠ EDGE-001 · 张力 2 · 状态 open

建议修订：Assign a specific role for post-update conflict auditing.

研究者回应：add_safeguard / add_safeguard — （研究者未填写理由）

⚠ EDGE-002 · 张力 2 · 状态 open

建议修订：Assign a specific role for post-update conflict auditing.

研究者回应：consult_stakeholders / consult_stakeholders — （研究者未填写理由）

研究者回应见上每条张力下方的「研究者回应」。

Step 6 · 修订台账

AI：已记录修订意向：

ID	议题	状态	改后文本	Replay
REV-001	Model drift may go unnoticed without a named owner to pause distribution and redesign the human-review workflow.	No concise open	changed passage was recorded.	Revision saved; replay evidence was not available.
REV-002	Model drift may go unnoticed without a named owner to pause distribution and redesign the human-review workflow.	No concise open	changed passage was recorded.	Revision saved; replay evidence was not available.
REV-003	Model drift may go unnoticed without a named owner to pause distribution and redesign the human-review workflow.	No concise open	changed passage was recorded.	Revision saved; replay evidence was not available.

来源：Desktop/SafeBARS_StudentTest/study_bundle.json